

EVEGLYPHDESIGN · MIRROR 3 · EgD-HANDBOOK-M3

Lillian's Guide

Transformational program management for information and technology

The Sovereign practitioner canon

The peer-review canon. This mirror of the handbook is the room-agnostic version, offered so consultants working in Lillian Corvington's lineage have something durable to practise from. It carries no client, no engagement, and no room-specific vocabulary. Where a specific room is being met, a tuned revision is published on a separate mirror; the peer-review canon stays here so the method can be judged on its own merits. Part VI reads the current SAP AI landscape — the Joule family, Business AI, AI Foundation, Business Data Cloud with Datasphere as its knowledge core — and names the boundary where the sovereign practice composes with it.

Written in the lineage of Lillian. The handbook is offered without ceremony; the reader is free to copy it, mirror it, and pass it along. Watermark and hash make tampering legible without preventing use.

Version rev 8 · Issued 2026-08-17T02:48:17Z

Source SHA-256 bd37c2288356b61dbb2cd4daa9a77e7b67a951825c5cadf662c2bc02603fec9e

Key ID EgD-KEY-2026-07 · Companion mirrors Mirror-4 Lillian's Guide (Group Platforms revision)

Transformational program management for information and technology

The Sovereign practitioner canon

Foreword. Something has changed in the way enterprises hold their information and technology function, and it changed more quickly than most of the industry has noticed. What used to be a service centre, carrying depreciation and operating expense, is coming to be held on the same terms as any other capital asset a firm might own. The systems of record, the data platforms, the models that reason over them, and the operating rituals around them — when they are built with provenance and portability and a rate of appreciation that can be measured — behave, on the ledger and in the boardroom, like the appreciating property they in fact are. The return on that property lands, as with any well-placed capital, in the parts of the business that are not IT: in receivables, in inventory, in cycle time, in the confidence of a regulator.

This is a handbook for practitioners who have decided to work in that way. It is written in the lineage of Lillian, who is the Sovereign of this practice — the word used the way it was designed to be used, as a piece of the design language of EVE Glyph Design and part of its copyright. The reader is not a bystander in that arrangement. The reader is the ruler of their sovereign data. Those are two different offices, and neither is a metaphor. The handbook is offered without ceremony; the reader is free to copy it, mirror it, and pass it along.

Part 0 — EVE, and why Lillian sits at the centre

EVE stands for Enterprise Value Engineering. It is the econometric utility ledger — the extension surface built on top of the ledger structure that most enterprises have already adopted, SAP's ACDOCA and its master-data companions — reaching outwards through Datasphere and any downstream structure the customer chooses to extend into.

EVE is pro-SAP. SAP is the foundation. There are commentators who argue that SAP is no longer relevant in the age of generative AI and sovereign data; we hold the opposite position, and hold it more strongly than the marketing does. The reason the work in this handbook is possible at all — attribution of profitability to the people who made it happen, ESG signals recorded at posting time, sovereign-AI memory over the customer's own data, industry roll-ups that actually reconcile — is that the customer already owns a full ledger substrate in ACDOCA and Datasphere. Nothing else in the market gives an enterprise that starting point.

Customer success is SAP success. That is SAP's own frame, and it is the right one. Our task is to help SAP's customers get more out of the capital investment they have already made, in SAP and in whatever else serves them, so that the value shows up on the customer's ledgers where it

belongs. Once the substrate is understood as the customer's own capital, the customer decides freely how the surrounding services are provided — by SAP, by one of the other capable firms, or by their own team.

Lillian is the practitioner in whose lineage the work is carried. EVE names the pattern; the handbook bears her name because the way she practises is what makes the pattern land.

Part I — What is actually the case

Four observations. Each is the kind of thing that becomes obvious once seen, and difficult to unsee.

Information and technology has become a capital asset. The line items that used to be treated as operating expense — data pipelines, reference models, evaluation harnesses, and the disciplined rituals around them — behave, when they are built with provenance and portability, like assets that appreciate. Held that way, they show up on the balance sheet the way other durable assets do, and they are managed accordingly.

The return on that asset is realised elsewhere. A well-structured data spine tends to reduce working capital in finance, shorten cycle time in operations, retain customers in commercial, and calm the disclosure conversation in legal. When the return is booked only inside the IT column, the asset is under-recognised, sometimes by an order of magnitude. Naming where the return actually lands, and marking it there, is a large part of what a good programme manager does.

The asset appreciates when it is owner-led. Rented systems and un-portable data appreciate on someone else's timetable, which is another way of saying they don't appreciate for the owner at all. When the customer holds the model of their own enterprise, the reference data, the evaluation set, and a tested exit, the asset compounds in the customer's hands. The consultant's role, when this is understood, becomes stewardship rather than tenancy — which is a healthier arrangement for both sides of the table.

The practice is portable; the practitioner is not. The methods in this handbook can be written down, and are. Written down, they can be studied, and are meant to be. But the judgement to know when to apply which method, and when to hold back and listen, comes from the person doing the work. That is why the handbook carries a name. It is not a claim about uniqueness; it is a note about where the practice actually lives.

Part II — How the work is done

There is a small handful of habits that, taken together, describe most of what a programme manager working in this lineage will do in the course of an ordinary week. They are habits, not

rules, and they are offered so a practitioner can recognise them when the situation calls for them.

Naming the appreciating asset before the charter. A programme that cannot name, in one line, the specific IT balance-sheet item it is here to appreciate is not yet ready to start. The naming exercise is short — an afternoon, sometimes an hour — and it saves months of drift later.

Marking the return where it lands. For every million a programme spends, there is a compensating figure that will show up somewhere else in the business, and it is worth taking the time to say where before the work begins. Receivables, inventory, cost-to-serve, risk-weighted capital. Published on day one, that map becomes the honest scoreboard the programme is judged against.

Owner-led custody, from the first drawing. Data, keys, evaluation sets, and orchestration definitions belong in the customer's repository, under the customer's identity provider, from the first architectural sketch. Designing this in later is much harder than designing it in from the start.

Keeping the mirror separate from the analysis. The system of record is captured faithfully; the analytic overlay lives alongside it, and is labelled as such. A programme that quietly rewrites the source is one that has failed a governance test no steering committee will catch until it is too late.

Working from the cheapest source that answers. Session context, memory, the wiki, the repository, and then, only if those have not produced what is needed, a targeted search. The programme's burn rate is a metric like any other, reported alongside schedule and scope.

Treating provenance as a delivery. Every artifact carries its hash, its timestamp, its author, its key ID, and its source lineage. Without those, an artifact is a draft, however polished it looks on the page.

Interrupting the sponsor only at the boundary of expensive spend. Free and cheap actions proceed. Expensive actions — new vendors, new environments, generation runs, capital commitments — pause for a written justification and a cheaper alternative. Sponsors appreciate not being pinged over nothing.

Making the effort visible. A slow answer that took real work reads as diligence; a slow answer that took no work reads as billing. The distinction is worth making legible in the moment, so the sponsor is not left to guess.

Two exits, tested. Every capitalised system has a documented, rehearsed exit to at least two alternatives. A customer who could walk away on a Tuesday morning and still operate on Wednesday is holding an asset in the full sense of the word.

Publishing the method, keeping the discretion. The method is public so it can be judged and improved. The customer's specifics stay private. That is the arrangement.

Part III — What "Sovereign" means, precisely

Sovereign, in this handbook, is a technical word. It describes a specific arrangement that can be checked against a customer's environment in an afternoon, without argument.

- Custody. The data, the models, the keys, and the orchestration definitions are stored where the customer controls them, under identities the customer provisions.
- Authority. Approvals, promotions, and material changes are signed by customer staff. Consultants recommend; owners decide.
- Exit. A tested route to at least two alternative platforms or vendors exists for every capitalised layer. Exit is rehearsed, not merely written.
- Return. The financial benefit of the appreciation is booked outside the IT function, on the ledgers where it is actually realised, and reviewed on the same cadence as any other capital asset.

A programme that satisfies all four is Sovereign in this sense. A programme that satisfies three is a programme with a specific vulnerability worth naming and, if the customer wishes, closing. A programme that satisfies fewer is a different kind of arrangement, and is best called by its ordinary name.

Part IV — The AI impact adaptations

The operating picture in Parts I–III is written so that a programme designed today does not carry the design decisions of a programme written two years ago. Three named adaptations sit on top of that picture, and each of them is a matter of what does not need to be built any more, so a practitioner can leave energy for what does. They are additions to the diagram, not replacements for it. The two-worlds structure, the ACDOCA spine, and the Decision Intelligence outcome are unchanged; the discovery scope defined during a current engagement is written knowing what the tools of today can now do.

OpEx efficiency, held inside the existing ERP estate. The operating-expense case for the transformation — process-cycle reduction, exception-handling reduction, working-capital release, licensing rationalisation — is calculated inside the existing ERP estate, on its own instruments, using its own reporting. That is the number the customer's finance function already trusts. The practitioner's contribution is not to invent a parallel OpEx measurement; it is to make the existing one legible against the programme's benefits map. This keeps the operating case defensible inside the enterprise after the engagement is over.

CapEx strategy, held inside the sovereign domain. The larger, longer-horizon capital strategy — what the transformation is worth to the ownership group over its planning horizon — is calculated inside the sovereign structure. The extended ledger holds the reference numbers; the PMO tool produces the case. The CapEx claim is a customer claim, made on the customer's own

instruments, defensible by the customer's own people.

PMO and business-case quantification, based on source data. Neither number — OpEx or CapEx — is estimated from a benchmark. Both are quantified from the enterprise's own source data. That is the discipline the sovereign structure exists to enable, and it is what makes the resulting numbers speakable inside the company after the engagement is over.

The three adaptations sit on the picture as an addition, not a redraw. What changes is that the discovery scope defined during a current engagement is written knowing what the tools of today can now do, rather than what the tools of two years ago could do.

Part V — The Discovery AI tool

The assessment layer is not a slide deck and it is not a spreadsheet. It is a small AI tool that runs inside the sovereign structure on the customer's side of the picture. It has four functions, in the order they run.

1. Analyse source sovereign data wherever it is today. Standard extractors against the ACDOCA layer per operating unit, or reads against an existing data lake if one is already in place. The tool goes to the data; the data does not move to the tool.
2. Ingest all existing third-party vendor materials into an AI-efficient, indexed, structured repository. Every prior assessment, every integrator deliverable, every scoping document, every stakeholder interview transcript, held once, indexed for retrieval, versioned. The repository is a residual asset in its own right — the customer keeps it after the engagement.
3. Identify quantitative economic benefits, prioritising technical leakage initially, and simultaneously consider the future-state primary technical barriers of the entire estate. Technical leakage is the first pass because it is the least contested and the fastest to quantify. The future-state barriers are held alongside so the leakage work does not accidentally invest in a pattern the future state will retire.
4. Identify early value-capture targets and begin feeding Decision Intelligence worklists to existing operational systems. The output is not a report. It is worklists that land in the operational systems the customer's people already use, so the value capture begins during the assessment rather than after it.

The tool is not a product. It is the assessment layer of the operating picture, named as a thing so it can be built, measured, and handed over. When the engagement ends, the tool stays inside the customer and continues to run.

Part VI — The SAP AI landscape, and where the sovereign practice sits on it

This handbook is written for practitioners running programmes on the SAP estate. It would be a defect of the peer-review canon to describe an assessment method, an AI adaptation tier, and a discovery tool without naming, in the same terms the room already uses, how the sovereign practice composes with the SAP AI stack that has been in place around the customer for years.

The section below names the current SAP AI construction, comments briefly on why the names keep changing, and states the boundary between what SAP owns by definition and what the sovereign practice adds. It also promotes SAP Datasphere from a target-state fabric to a discovery instrument, because that promotion is a load-bearing part of how the sovereign cadence recovers cost against the standard cadence.

1. What SAP currently names, and how the naming is settling

SAP's current top-level construction, as SAP itself names it on the customer-facing product pages today:

- Joule is the umbrella brand for SAP's AI experience.
- Joule Assistants are the embedded, grounded assistants inside the individual applications. Published by business function: Cash and Treasury, Buying, Business Network, Campaign, Career and Talent Development, Category Management, Compensation, Content, Core HR, Procurement Contract, Recurring Receivables, Workforce Upskilling, and a growing catalogue. Joule base activates for free on existing cloud subscriptions.
- Joule Agents are non-deterministic workflow agents. They choose from Joule skills, other agents, and third-party applications; act, reflect, and reason on results; and are grounded on SAP Knowledge Graph and SAP Business Data Cloud.
- Joule Studio is the build lane for enterprise-wide agentic AI — the surface for the customer's own agents.
- Joule for Developers sits alongside Joule Studio for developer-productivity workloads. Named in the Bosch Digital reference alongside SAP AI Core and ABAP Cloud.
- SAP Business AI is the enterprise-wide brand covering the embedded predictive models and closed-loop feedback inside individual applications.
- SAP AI Core persists as the runtime layer beneath the Joule experience, still explicitly named in current customer stories.
- SAP AI Foundation is where the generative-AI hub, RAG plumbing, prompt-management, and grounding services are described in SAP's platform material. Its landing page is currently

intermittently reachable, which is a small but useful piece of evidence that SAP is still settling how it names this layer.

The data and reasoning layer beneath all of this:

- SAP Business Data Cloud is the current umbrella architecture. In SAP's own words: SAP Datasphere is a key component in SAP Business Data Cloud, serving as the knowledge core that unifies semantics, data products, and modeling.
- SAP Datasphere is the semantic and modeling core inside that fabric. It provides data integration, cataloging, semantic modeling, data warehousing, data-product management, discovery and activation, and virtualization.
- SAP Knowledge Graph is the graph structure over the SAP estate that Joule Agents ground on.
- SAP HANA Cloud is the runtime for the ledger, the vector and graph engines, and the reasoning surface beneath the Knowledge Graph.
- SAP Databricks is SAP's own naming for the data-science and AI-processing lane on the fabric.
- SAP Analytics Cloud is the visualisation and planning surface on the fabric.
- Data Marketplace and Catalog are the SAP-managed data-product surfaces reachable through Datasphere.

The generative-AI hub composes across external LLMs (SAP has described compositions across Anthropic, OpenAI, Mistral, Meta, and hyperscaler-hosted models in prior TechEd and Sapphire material). The current customer-facing pages do not list those providers by name, and the composition changes.

2. The renaming trail, and what it means for a programme

In the past three years SAP has moved through, roughly in order: SAP AI Business Services → SAP AI Core → SAP AI Launchpad → SAP AI Foundation → SAP Business AI → the Joule family (Joule, Joule Assistants, Joule Agents, Joule Studio, Joule for Developers) → SAP Business Data Cloud absorbing SAP Datasphere as its knowledge core.

None of these names is wrong. Each renaming closes an older packaging problem and opens a newer one. The lesson for a transformational programme is not to attach roadmap milestones to the current names; it is to attach them to the underlying capabilities — grounded chat, agent authoring, embedded prediction, generative hub, semantic modeling, data-product publication, knowledge-graph reasoning — and let the names update themselves in the programme documents when SAP updates them next.

The sovereign practice takes the opposite posture: it stabilises what it names, and names it in the customer's terms rather than the vendor's. The two postures are complementary. The vendor is entitled to iterate the packaging of its offering; the customer is entitled to a lane of the estate that

keeps its names, its ledger structure, and its rituals across those iterations.

3. Where the sovereign practice sits on the map

The sovereign practice is not a competitor to any of the components above. It sits above the SAP ledger and reaches outward through the business data fabric. The boundary line is drawable, and it is worth drawing precisely.

- Joule Assistants and Joule Agents run inside the SAP-owned scenarios. Their grounding is the suite's own data and process context. They are the correct instrument for anything living inside the SAP process boundary. The sovereign practice does not replace them.
- Joule Studio is the correct place to author SAP-native agents. For anything whose entire behaviour lives inside the SAP process boundary, that is where it belongs.
- The sovereign practice begins where the SAP process boundary ends. Cross-suite econometric reasoning, capital-appreciation of the information-and-technology asset on the balance sheet, cross-brand signal that has to cross SAP scenarios to be legible, sovereignty over models and rituals that must not leave the customer's ledger, audit lanes that must hold when a partner LLM composition changes underneath. That is the layer this handbook calls EVE.
- The two lanes are composable, not adjacent. Joule Agents can call the sovereign layer for context they do not natively hold; the sovereign layer can query Joule Assistants and Agents for scenario-specific answers. Neither displaces the other. The interface between them is the business data fabric.

A useful one-line statement of the position: where Joule stops, EVE begins. Joule is not stopped by weakness; it is stopped by scope. The sovereign lane holds what the scope leaves out.

4. SAP Datasphere before conversion — a discovery instrument, not a target-state fabric

The standard S/4HANA conversion cadence deploys Datasphere late — after conversion, as the fabric that sits on top of the new target state. That is a reasonable target-state posture. It is a poor discovery posture, because it forces the discovery phase to be run against interviews, diagrams, and the recollected shape of the current estate rather than against the current estate itself.

The sovereign cadence pre-deploys Datasphere connectivity to the current estate as a discovery instrument. ECC, satellite ERPs, master-data hubs, extract lakes, warehouse marts. Datasphere is not asked to run the current estate; it is asked to observe it — to catalog the actual sources, surface the actual quality signal, expose the actual master-data drift, and name the data products the business has already built informally on top of the SAP core.

Four consequences follow:

1. Discovery moves from memory to measurement. Requirements are gathered against the estate as it actually behaves, not as it is recalled. Interviews still happen, but they are now interpreting measurement rather than substituting for it.
2. Data quality is addressed before conversion, not during it. The poorest-quality feeds are named early; master-data reconciliation is scheduled ahead of cutover; harmonisation logic is drafted against real distributions rather than field lists. Cutover is not the place to discover data problems.
3. Analytics is not blocked by the conversion cadence. The fabric is already in place upstream of the S/4HANA move; reporting, planning, and Joule-Agent grounding do not have to wait for target-state cutover to be trustworthy.
4. The Discovery AI tool has a real substrate. The four functions named in Part V are much cheaper to run when Datasphere has already surfaced the estate. The chain — Datasphere for observation, Discovery AI for reasoning over that observation, EVE for stable custody of the reasoning — is the sovereign cadence's structural claim on cost recovery against the standard cadence.

The pre-deployment is a chargeable piece of platform work. Its return is measured not in analytics deliverables but in conversions arriving on estates that have already been named, measured, and cleaned up in the year prior — instead of estates being named for the first time in the middle of a cutover.

In plain terms: see what you have before you decide what to build.

Datasphere-as-discovery-instrument is how that sentence lands as a programme decision, in the vocabulary already in use in the room.

Part VII — PAIX, and the wider work

Pro bono humanitatis.

PAIX is the French word for peace, and the name we use for the wider work — the part of the practice that sits above any individual client engagement. Most of what a practitioner in this lineage does in an ordinary year happens here. Two habits are worth describing.

Going to where people are, rather than waiting for them to become customers. The tools that used to take a room full of confusion and turn it into a clear set of requirements lived, for a long time, inside expensive enterprise stacks. To use them at all, one needed a customer, a licence, a project. That is no longer quite true. The tools now within reach of anyone with a browser can do a great deal of what those stacks did at the requirements stage. So there is nothing that keeps a practitioner from taking those tools into the communities they live in and helping the people there use them well — whether or not those communities will ever be customers of the practice. When a larger commercial system does eventually arrive, the requirements have been thought through, the

vocabulary has been agreed, and the people who will run it already know what they are running. That is the work; it does not need to be called anything else.

Staying with people after the engagement ends. Client relationships come and go on their own schedule. The people met inside those relationships remain, in a quieter way, part of the practice. When a career needs a hand at recovery, we help. When a call comes in with something difficult, long after the original engagement is closed, we take the call. This is why the referral network is offered freely and stays offered, and why a first consultation can honestly say that anything useful said in the room is the client's to keep.

The habit, taken together, is a habit of adding — a method, a tool now within reach, a name to call, a competent hand held out at the right moment. Where we can add, we add. Where we cannot, we say so, and we point to someone who can.

Author's note — a calibration line

Every honest practitioner keeps a self-correction near the desk, because the drift from outcome into steps is the drift that every method in every discipline is vulnerable to. The correction preserved here is the one this handbook is measured against.

You are still describing the steps to do rather than the expected outcome — i.e. why and how we understand what sovereign data means. Begin by decoding the technical leakage. Discovery worklists to existing operational systems. What is existing? What is better, what is different, what is missing?

Read the practice above against that correction. The outcome is what sovereign data means, made speakable inside the enterprise by the enterprise's own people, on the enterprise's own instruments. The entry point is decoding technical leakage — the least contested, fastest quantifiable, and most defensible starting point for the assessment. The delivery is worklists into existing operational systems — not new dashboards, not new meetings, not new committees.

Any future revision of this handbook that drifts from that outcome, that entry point, or that delivery is a drift. The self-correction is the calibration line.

Colophon

Written in the lineage of Lillian, for practitioners doing the work of Enterprise Value Engineering — the econometric utility ledger extending customer-owned SAP and Datasphere substrates. Watermarked and content-hashed on export. The reader is free to copy, mirror, and redistribute this document; the watermark makes tampering legible without preventing use.

Pour le bien-être du peuple.

EVEglyphDesign