

EVEGLYPHDESIGN · MIRROR 4 · EgD-HANDBOOK-M4

# Lillian's Guide

## Transformational program management for information & technology

The Sovereign practitioner canon

For this reader. This mirror of the handbook is prepared for a Group Platforms conversation inside a multi-brand enterprise that has already made SAP its foundation, is preparing the next wave of its S/4HANA journey through a business-led, clean-core, satellites-and-BTP posture, and is currently completing discovery deliverables in advance of solution design. The vocabulary in the pages that follow is chosen to sit inside that posture. Part VI reads the current SAP AI landscape — Joule, Joule Studio, Joule Agents, Business AI, AI Foundation, Business Data Cloud with Datasphere as its knowledge core — and states where the sovereign practice composes with it and where it does not.

Written in the lineage of Lillian. The handbook is offered without ceremony; the reader is free to copy it, mirror it, and pass it along. Watermark and hash make tampering legible without preventing use.

Version rev 11 · Issued 2026-08-17T02:48:16Z

Source SHA-256 abf743b639a4ab03bf1adadc59cb8e68dc1c448493297a42f647d997303497d7

Key ID EgD-KEY-2026-07 · Companion Mirror-3 Lillian's Guide (peer-review canon)

# Transformational program management for information & technology

---

## The Sovereign practitioner canon

For this reader. This mirror of the handbook is prepared for a Group Platforms conversation inside a multi-brand enterprise that has already made SAP its foundation, is preparing the next wave of its S/4HANA journey through a business-led, clean-core, satellites-and-BTP posture, and is currently completing discovery deliverables in advance of solution design. The vocabulary in the pages that follow is chosen to sit inside that posture. Nothing in the handbook asks the reader to translate; the practice is written in the terms already in use in the room.

Foreword. Something has changed in the way enterprises hold their information and technology function, and it changed more quickly than most of the industry has noticed. What used to be a service centre, carrying depreciation and operating expense, is coming to be held on the same terms as any other capital asset a firm might own. The systems of record — an ERP core kept clean, the satellites that reach outward from it, the extensions that live on the platform layer, the data spine that binds them, the models that reason over that data, and the operating rituals around them — when they are built with provenance and portability and a rate of appreciation that can be measured, behave, on the ledger and in the boardroom, like the appreciating property they in fact are. The return on that property lands, as with any well-placed capital, in the parts of the business that are not the IT function: in receivables, in inventory, in cycle time, in the confidence of a regulator, in the quality of a brand-to-brand signal that used to travel through email.

This is a handbook for practitioners who have decided to work in that way. It is written in the lineage of Lillian, who is the Sovereign of this practice — the word used the way it was designed to be used, as a piece of the design language of EVE Glyph Design and part of its copyright. The reader is not a bystander in that arrangement. The reader is the ruler of their sovereign data. Those are two different offices, and neither is a metaphor. The handbook is offered without ceremony; the reader is free to copy it, mirror it, and pass it along.

## Part 0 — EVE, and where it stands in a clean-core landscape

---

EVE stands for Enterprise Value Engineering. It is the econometric utility ledger — the extension surface built on top of the ledger structure that most enterprises have already adopted, SAP's ACDOCA and its master-data companions — reaching outwards through Datasphere and any downstream structure the customer chooses to extend into.

EVE is pro-SAP. SAP is the foundation. There are commentators who argue that SAP is no longer relevant in the age of generative AI and sovereign data; we hold the opposite position, and hold it more strongly than the marketing does. The reason the work in this handbook is possible at all — attribution of profitability to the people who made it happen, ESG signals recorded at posting time, sovereign-AI memory over the customer's own data, industry roll-ups that actually reconcile, brand-level margin walks that stand up to a group financial review — is that the customer already owns a full ledger substrate in ACDOCA and Datasphere. Nothing else in the market gives an enterprise that starting point.

EVE sits as a satellite, inside the clean-core envelope. It does not modify the ERP core. It does not compete with the group data platform. It runs on the extension platform the enterprise is already committed to — the SAP Business Technology Platform where extensions belong — and it contributes semantics upward to the group data platform through the same integration principles the enterprise's own architects have written. Where a Chief Data Officer is building a data-and-AI-augmented future for the enterprise, EVE is additive to that build, contributing the enterprise-value axis to the semantics the platform already carries. Where a Chief Platforms Officer is protecting a clean core, EVE is a satellite investment measurable against the programme's benefits case — not a modification of the core the programme is protecting.

Customer success is SAP success. That is SAP's own frame, and it is the right one. Our task is to help SAP's customers get more out of the capital investment they have already made, in SAP and in whatever else serves them, so that the value shows up on the customer's ledgers where it belongs. Once the substrate is understood as the customer's own capital, the customer decides freely how the surrounding services are provided — by SAP, by one of the other capable firms, or by their own team.

Lillian is the practitioner in whose lineage the work is carried. EVE names the pattern; the handbook bears her name because the way she practises is what makes the pattern land.

## Part I — What is actually the case

Four observations. Each is the kind of thing that becomes obvious once seen, and difficult to unsee.

Information and technology has become a capital asset. The line items that used to be treated as operating expense — data pipelines, reference models, evaluation harnesses, and the disciplined rituals around them — behave, when they are built with provenance and portability, like assets that appreciate. Held that way, they show up on the balance sheet the way other durable assets do, and they are managed accordingly. Inside a business-led ERP transformation, the appreciating asset is the simplified and aligned enterprise landscape itself — the coherent-and-future-proof landscape that a Group Platforms function is trying to leave behind for the next decade.

The return on that asset is realised elsewhere. A well-structured data spine tends to reduce working capital in finance, shorten cycle time in operations, retain customers in commercial, and calm the disclosure conversation in legal. In a multi-entity, multi-brand enterprise it also shortens the distance between what an operating brand knows about its own performance and what the group can see at consolidation, which is a distinct benefit worth naming on its own. When the return is booked only inside the IT column, the asset is under-recognised, sometimes by an order of magnitude. Naming where the return actually lands, and marking it there, is a large part of what a good programme manager does.

The asset appreciates when it is owner-led. Rented systems and un-portable data appreciate on someone else's timetable, which is another way of saying they don't appreciate for the owner at all. When the customer holds the model of their own enterprise, the reference data, the evaluation set, and a tested exit, the asset compounds in the customer's hands. The consultant's role, when this is understood, becomes stewardship rather than tenancy — which is a healthier arrangement for both sides of the table, and specifically the arrangement a business-led programme is designed to produce.

The practice is portable; the practitioner is not. The methods in this handbook can be written down, and are. Written down, they can be studied, and are meant to be. But the judgement to know when to apply which method, and when to hold back and listen, comes from the person doing the work. That is why the handbook carries a name. It is not a claim about uniqueness; it is a note about where the practice actually lives.

## Part II — How the work is done

There is a small handful of habits that, taken together, describe most of what a programme manager working in this lineage will do in the course of an ordinary week. They are habits, not rules, and they are offered so a practitioner can recognise them when the situation calls for them.

Naming the appreciating asset before the charter. A programme that cannot name, in one line, the specific balance-sheet item it is here to appreciate is not yet ready to start. The naming exercise is short — an afternoon, sometimes an hour — and it saves months of drift later. In a discovery-into-solution-design flow, the naming exercise is the first artifact of solution design, not a preamble to it.

Marking the return where it lands. For every million a programme spends, there is a compensating figure that will show up somewhere else in the business, and it is worth taking the time to say where before the work begins. Receivables, inventory, cost-to-serve, risk-weighted capital, brand-level margin. Published on day one, that map becomes the honest scoreboard the programme is judged against — and, at a group level, the same map is what a Technology CFO uses to defend the platform's unit economics against the operating businesses whose money it is spending.

Owner-led custody, from the first drawing. Data, keys, evaluation sets, and orchestration definitions belong in the customer's repository, under the customer's identity provider, from the first architectural sketch. Designing this in later is much harder than designing it in from the start. Inside a clean-core doctrine, owner-led custody is what makes the satellite defensible: the enterprise can extend, contract, or replace the satellite without touching the core it is protecting.

Keeping the satellite outside the core. The ERP core is kept clean. Extensions live on the extension platform the enterprise has committed to, with clear boundaries and named integration points. An EVE deployment that would require a modification of the core is, on this point of principle, not the right EVE deployment for the enterprise; the practice adjusts the satellite to keep the core intact.

Keeping the mirror separate from the analysis. The system of record is captured faithfully; the analytic overlay lives alongside it, and is labelled as such. A programme that quietly rewrites the source is one that has failed a governance test no steering committee will catch until it is too late.

Working from the cheapest source that answers. Session context, memory, the wiki, the repository, and then, only if those have not produced what is needed, a targeted search. The programme's burn rate is a metric like any other, reported alongside schedule and scope.

Treating provenance as a delivery. Every artifact carries its hash, its timestamp, its author, its key ID, and its source lineage. Without those, an artifact is a draft, however polished it looks on the page.

Interrupting the sponsor only at the boundary of expensive spend. Free and cheap actions proceed. Expensive actions — new vendors, new environments, generation runs, capital commitments — pause for a written justification and a cheaper alternative. Sponsors appreciate not being pinged over nothing.

Making the effort visible. A slow answer that took real work reads as diligence; a slow answer that took no work reads as billing. The distinction is worth making legible in the moment, so the sponsor is not left to guess.

Two exits, tested. Every capitalised system has a documented, rehearsed exit to at least two alternatives. A customer who could walk away on a Tuesday morning and still operate on Wednesday is holding an asset in the full sense of the word.

Publishing the method, keeping the discretion. The method is public so it can be judged and improved. The customer's specifics stay private. That is the arrangement.

## Part III — What "Sovereign" means, precisely

Sovereign, in this handbook, is a technical word. It describes a specific arrangement that can be checked against a customer's environment in an afternoon, without argument.

- Custody. The data, the models, the keys, and the orchestration definitions are stored where the customer controls them, under identities the customer provisions.
- Authority. Approvals, promotions, and material changes are signed by customer staff. Consultants recommend; owners decide.
- Exit. A tested route to at least two alternative platforms or vendors exists for every capitalised layer. Exit is rehearsed, not merely written.
- Return. The financial benefit of the appreciation is booked outside the IT function, on the ledgers where it is actually realised, and reviewed on the same cadence as any other capital asset.

A programme that satisfies all four is Sovereign in this sense. A programme that satisfies three is a programme with a specific vulnerability worth naming and, if the customer wishes, closing. A programme that satisfies fewer is a different kind of arrangement, and is best called by its ordinary name.

Inside a business-led ERP transformation running on a clean-core doctrine, the Sovereign checklist is compatible with the programme's own principles. Custody and authority are what "business-led" means when the phrase is taken seriously; exit is what a coherent-and-future-proof landscape is supposed to preserve for its owner; return is what the benefits case is measuring in the first place. A programme that passes its own principles will already pass most of this checklist; the practice's contribution is the discipline of saying so, in writing, at solution-design time, before the shape of the answer forecloses the option.

## Part IV — The AI impact adaptations

The picture of the practice has always held two worlds. On one side is the customer's sovereign structure — the ledger substrate, the reference data, the evaluation set, the operating rituals — held inside the customer's own environment and never rented back. On the other side is the operational reality of the enterprise as it runs today, on the systems it already runs. The transformation is the act of holding both worlds in view at once, and letting the sovereign structure absorb what belongs on it while leaving what belongs in the operating systems where it is.

What has changed since the last generation of programmes were designed is the toolset available to run the assessment layer that sits between those two worlds. Three adaptations follow from that change, and they are named here so a programme written today does not carry the design decisions of a programme written two years ago.

OpEx efficiency, held inside SAP. The operational efficiency case that AI now makes reachable is measured inside the existing SAP estate. Nothing about the OpEx work moves data out of the current systems or asks the enterprise to buy a new tool. The instrumentation lives on the axis, not on either side of it.

CapEx strategy, held inside the sovereign domain. The larger, longer-horizon capital strategy — what the transformation is worth to the ownership group over its planning horizon — is calculated inside the sovereign structure. The extended ledger holds the reference numbers; the PMO tool produces the case. The CapEx claim is a customer claim, made on the customer's own instruments, defensible by the customer's own people.

PMO and business-case quantification, based on source data. Neither number — OpEx or CapEx — is estimated from a benchmark. Both are quantified from the enterprise's own source data. That is the discipline the sovereign structure exists to enable, and it is what makes the resulting numbers speakable inside the company after the engagement is over.

The three adaptations sit on the picture as an addition, not a redraw. The two-worlds structure, the ACDOCA spine, and the Decision Intelligence Centre of Excellence outcome do not change. What changes is that the discovery scope defined during a current engagement is written knowing what the tools of today can now do, rather than what the tools of two years ago could do.

## Part V — The Discovery AI tool

The assessment layer is not a slide deck and it is not a spreadsheet. It is a small AI tool that runs inside the sovereign structure on the customer's side of the picture. It has four functions, in the order they run.

1. Analyse source sovereign data wherever it is today. Standard extractors against the ACDOCA layer per operating brand, or reads against an existing data lake if one is already in place. The tool goes to the data; the data does not move to the tool.
2. Ingest all existing third-party vendor materials into an AI-efficient, indexed, structured repository. Every prior assessment, every integrator deliverable, every scoping document, every stakeholder interview transcript, held once, indexed for retrieval, versioned. The repository is a residual asset in its own right — the customer keeps it after the engagement.
3. Identify quantitative economic benefits, prioritising technical leakage initially, and simultaneously consider the future-state primary technical barriers of the entire estate. Technical leakage is the first pass because it is the least contested and the fastest to quantify. The future-state barriers are held alongside so the leakage work does not accidentally invest in a pattern the future state will retire.
4. Identify early value-capture targets and begin feeding Decision Intelligence worklists to existing operational systems. The output is not a report. It is worklists that land in the operational systems the customer's people already use, so the value capture begins during the assessment rather than after it.

The tool is not a product. It is the assessment layer of the operating diagram, named as a thing so it can be built, measured, and handed over. When the engagement ends, the tool stays inside the



customer and continues to run.

## Part VI — Inside the SAP AI landscape

The reader has been pumped full of SAP's AI story since the first cloud briefing, and rightly so. It is a serious offering and the room already runs on the ledger underneath it. What follows is not a review of that offering; it is a map of it, so this handbook can name where the sovereign practice sits on the same board, and where the two lanes complement each other cleanly.

### What SAP is currently naming

SAP has settled, for the moment, on the following top-level construction. The names are what they are on the customer-facing product pages today; they have moved before and will move again.

- Joule is the umbrella brand for SAP's AI experience. Everything below sits under it.
- Joule Assistants are the grounded assistants embedded in the individual applications, published by business function: Cash and Treasury, Buying, Category Management, Core HR, Compensation, Recurring Receivables, Workforce Upskilling, and a growing catalogue. Joule base activates for free on existing cloud subscriptions.
- Joule Agents are non-deterministic workflow agents. They choose from Joule skills, other agents, and third-party applications, act, reflect, and reason. Grounded on SAP Knowledge Graph and SAP Business Data Cloud.
- Joule Studio is the build lane for enterprise-wide agentic AI — where the customer authors its own agents inside SAP's runtime.
- Joule for Developers and SAP AI Core persist as the developer-productivity and runtime layers beneath the experience.
- SAP Business AI is the enterprise-wide brand covering embedded predictive models and closed-loop feedback inside individual applications.

Beneath the AI surface, on the data and reasoning layer:

- SAP Business Data Cloud is the current umbrella architecture. In SAP's own words: SAP Datasphere is a key component in SAP Business Data Cloud, serving as the knowledge core.
- SAP Datasphere provides data integration, cataloging, semantic modeling, data warehousing, data-product management, discovery and activation, and virtualization inside that fabric.
- SAP Knowledge Graph, SAP HANA Cloud, SAP Databricks, SAP Analytics Cloud, and the Data Marketplace complete the composition.
- The generative-AI hub inside SAP AI Foundation composes across external LLMs — SAP has described compositions across Anthropic, OpenAI, Mistral, Meta, and hyperscaler-hosted



models in TechEd and Sapphire material. The customer-facing pages do not currently list those providers by name, and the AI Foundation landing page is intermittently reachable — evidence that SAP is still settling how it names this layer.

## The renaming trail

In the past three years SAP has moved through: SAP AI Business Services → SAP AI Core → SAP AI Launchpad → SAP AI Foundation → SAP Business AI → the Joule family (Joule, Joule Assistants, Joule Agents, Joule Studio, Joule for Developers) → SAP Business Data Cloud absorbing SAP Datasphere as its knowledge core. None of these names is wrong. Each renaming closes an older packaging problem and opens a newer one.

The practical read for a Group Platforms roadmap: name what you deploy by the current name, expect it to be renamed once mid-programme, and attach roadmap milestones to the underlying capabilities (grounded chat, agent authoring, embedded prediction, generative hub, semantic modeling, data-product publication, knowledge-graph reasoning) rather than to the current badges.

## Where the sovereign practice sits — where Joule stops, EVE begins

The sovereign practice is not a competitor to any of the components above. It is what sits above the SAP ledger and reaches outward through the business data fabric. The boundary is worth drawing precisely, because in this kind of room the wrong framing turns a complementary layer into a threat.

- Joule Assistants and Joule Agents run inside the SAP-owned scenarios. Their grounding is the suite's own data and process context. They are the correct instrument for anything living inside the SAP process boundary. The sovereign practice does not touch that layer.
- Joule Studio is where the room builds its own SAP-native agents. For anything whose behaviour lives entirely inside the SAP process boundary, that is where it belongs.
- The sovereign practice begins where the SAP process boundary ends. Cross-suite econometric reasoning — capital-appreciation on the information-and-technology asset, cross-brand signal that has to cross scenarios to be legible, sovereignty over models and rituals that must not leave the customer's ledger, audit lanes that must hold when a partner LLM composition changes. That is the layer this handbook calls EVE.
- The two lanes are composable, not adjacent. Joule Agents can call the sovereign layer for context they do not natively hold; the sovereign layer can query Joule for scenario-specific answers. Neither displaces the other. The interface between them is the business data fabric.

One line that lands the position: where Joule stops, EVE begins. Joule is not stopped by weakness; it is stopped by scope. The sovereign lane carries the questions that fall outside Joule's scope by design.

## SAP Datasphere before conversion — a discovery instrument, not a target-state fabric

The standard S/4HANA cadence deploys Datasphere late — after conversion, as the fabric that sits on top of the new target state. That is a reasonable target-state posture. It is a poor discovery posture, because it forces the discovery phase to be run against interviews, diagrams, and the recollected shape of the current estate rather than against the current estate itself.

The sovereign cadence pre-deploys Datasphere connectivity to the current estate — ECC, satellite ERPs, master-data hubs, extract lakes, warehouse marts — as a discovery instrument. Datasphere is not being asked to run the current estate; it is being asked to observe it. To catalog the actual sources, surface the actual quality signal, expose the actual master-data drift, and name the data products the business has already built informally on top of the SAP core.

Four consequences follow. Discovery moves from memory to measurement. Data quality is addressed before conversion, not during it. Analytics is not blocked by the conversion cadence, because the fabric is already in place upstream of it. And the Discovery AI tool in Part V has a real substrate — Datasphere-collected observation — to reason over, so the chain of Datasphere-for-observation, Discovery-AI-for-reasoning, EVE-for-stable-custody becomes the sovereign cadence's structural claim on cost recovery against the standard cadence.

The pre-deployment is a chargeable piece of platform work. Its return is measured not in analytics deliverables but in conversions arriving on estates that have already been named, measured, and cleaned up in the year prior. In plain terms: see what you have before you decide what to build.

## Part VII — PAIX, and the wider work

Pro bono humanitatis.

PAIX is the French word for peace, and the name we use for the wider work — the part of the practice that sits above any individual client engagement. Most of what a practitioner in this lineage does in an ordinary year happens here. Two habits are worth describing.

Going to where people are, rather than waiting for them to become customers. The tools that used to take a room full of confusion and turn it into a clear set of requirements lived, for a long time, inside expensive enterprise stacks. To use them at all, one needed a customer, a licence, a project. That is no longer quite true. The tools now within reach of anyone with a browser can do a great deal of what those stacks did at the requirements stage. So there is nothing that keeps a practitioner from taking those tools into the communities they live in and helping the people there use them well — whether or not those communities will ever be customers of the practice. When a larger commercial system does eventually arrive, the requirements have been thought through, the vocabulary has been agreed, and the people who will run it already know what they are running. That is the work; it does not need to be called anything else.

Staying with people after the engagement ends. Client relationships come and go on their own schedule. The people met inside those relationships remain, in a quieter way, part of the practice. When a career needs a hand at recovery, we help. When a call comes in with something difficult, long after the original engagement is closed, we take the call. This is why the referral network is offered freely and stays offered, and why a first consultation can honestly say that anything useful said in the room is the client's to keep.

The habit, taken together, is a habit of adding — a method, a tool now within reach, a name to call, a competent hand held out at the right moment. Where we can add, we add. Where we cannot, we say so, and we point to someone who can.

## Author's note — a calibration line

Every honest practitioner keeps a self-correction near the desk, because the drift from outcome into steps is the drift that every method in every discipline is vulnerable to. The correction preserved here is the one this handbook is measured against.

You are still describing the steps to do rather than the expected outcome — i.e. why and how we understand what sovereign data means. Begin by decoding the technical leakage. Discovery worklists to existing operational systems. What is existing? What is better, what is different, what is missing?

Read the practice above against that correction. The outcome is what sovereign data means, made speakable inside the enterprise by the enterprise's own people, on the enterprise's own instruments. The entry point is decoding technical leakage — the least contested, fastest quantifiable, and most defensible starting point for the assessment. The delivery is worklists into existing operational systems — not new dashboards, not new meetings, not new committees.

Any future revision of this handbook that drifts from that outcome, that entry point, or that delivery is a drift. The self-correction is the calibration line.

## Colophon

Written in the lineage of Lillian, for practitioners doing the work of Enterprise Value Engineering — the econometric utility ledger extending customer-owned SAP and Datasphere substrates as a satellite on the extension platform, inside a clean-core doctrine. Watermarked and content-hashed on export. The reader is free to copy, mirror, and redistribute this document; the watermark makes tampering legible without preventing use.

Pour le bien-être du peuple.